

# การจำแนกการออกเสียง ภาษาไทยของชาวต่างชาติ

Thai Pronunciation of foreigners Classification

# Introduction

Marketeer

## มหาวิทยาลัยเอกชน เติบโตด้วยนักศึกษาต่างชาติ



2558

8,255 คน

2559

9,840 คน

2560

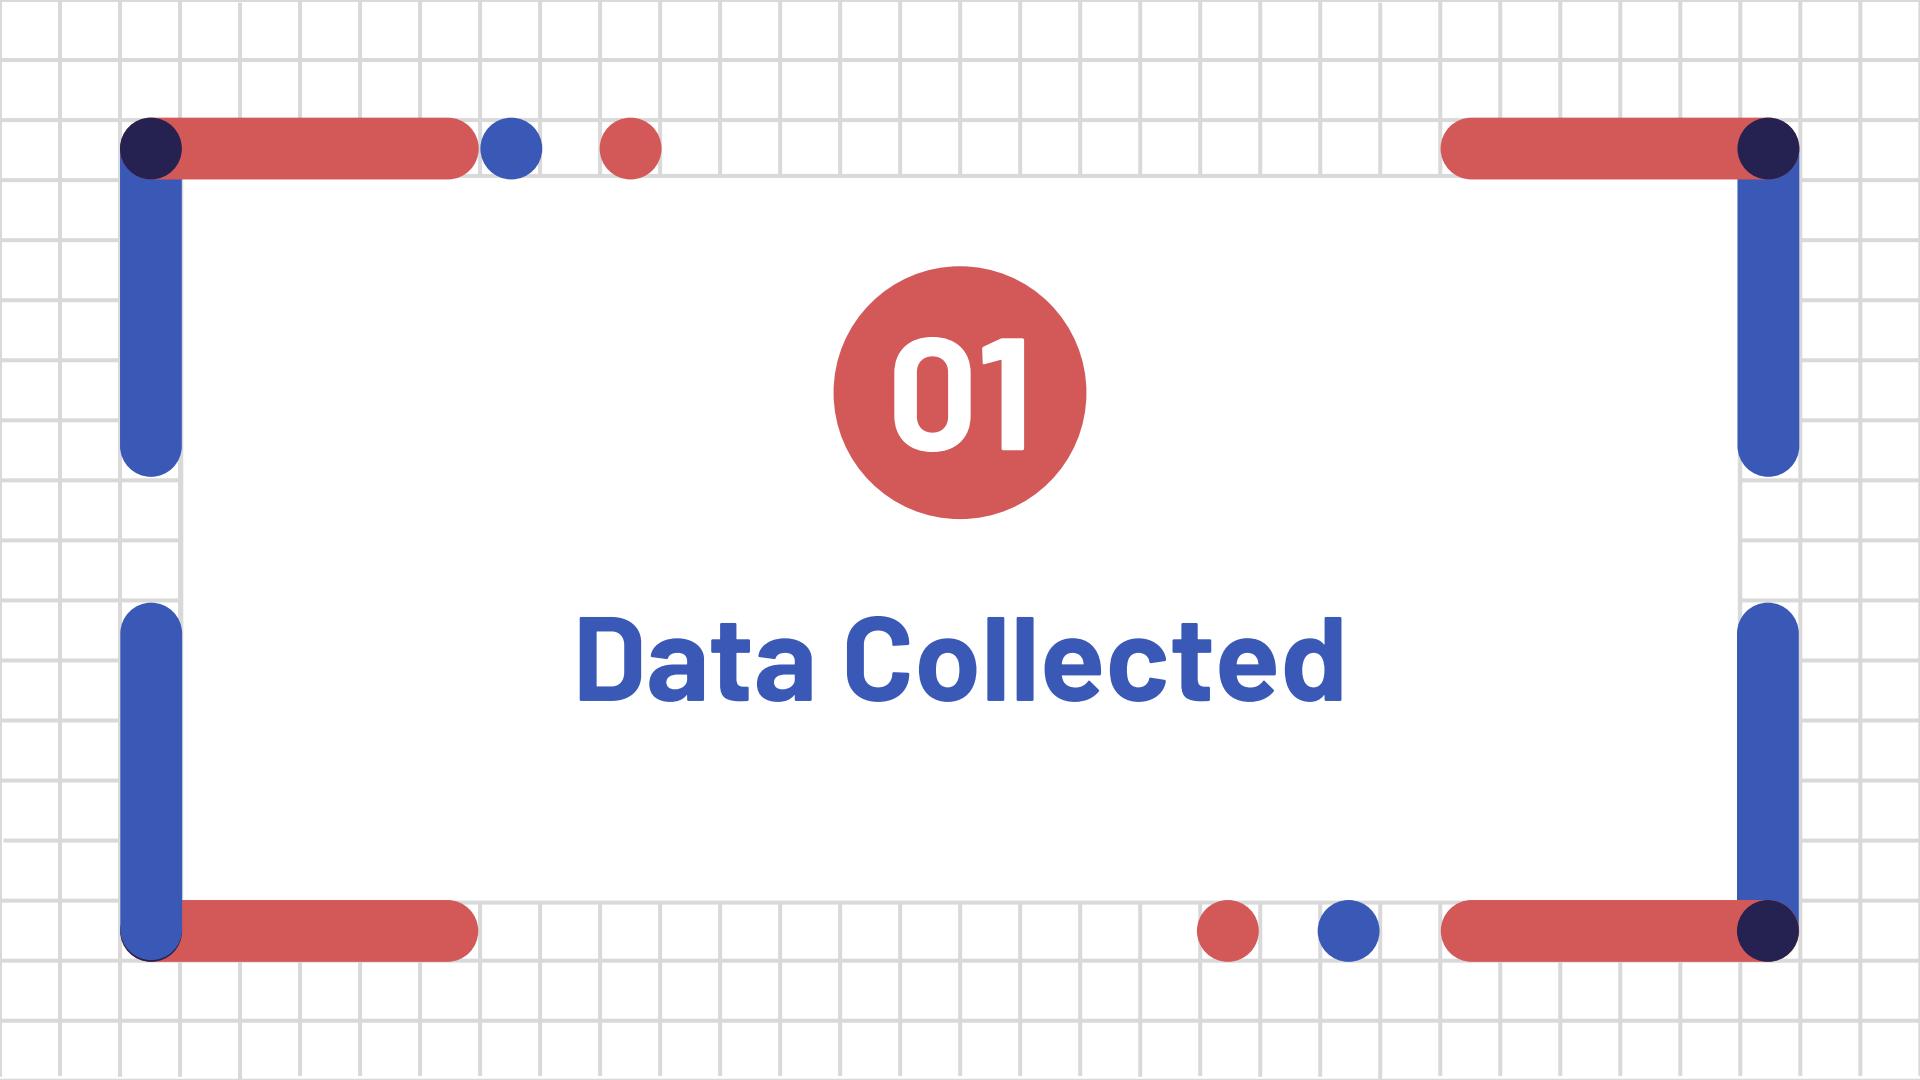
11,512 คน



\*นักศึกษาใหม่ประจำปีการศึกษา

ที่มา : สำนักงานคณะกรรมการอุดมศึกษา, วางแผนจากนักศึกษาต่างชาติที่เข้ามาในปีการศึกษา 2560





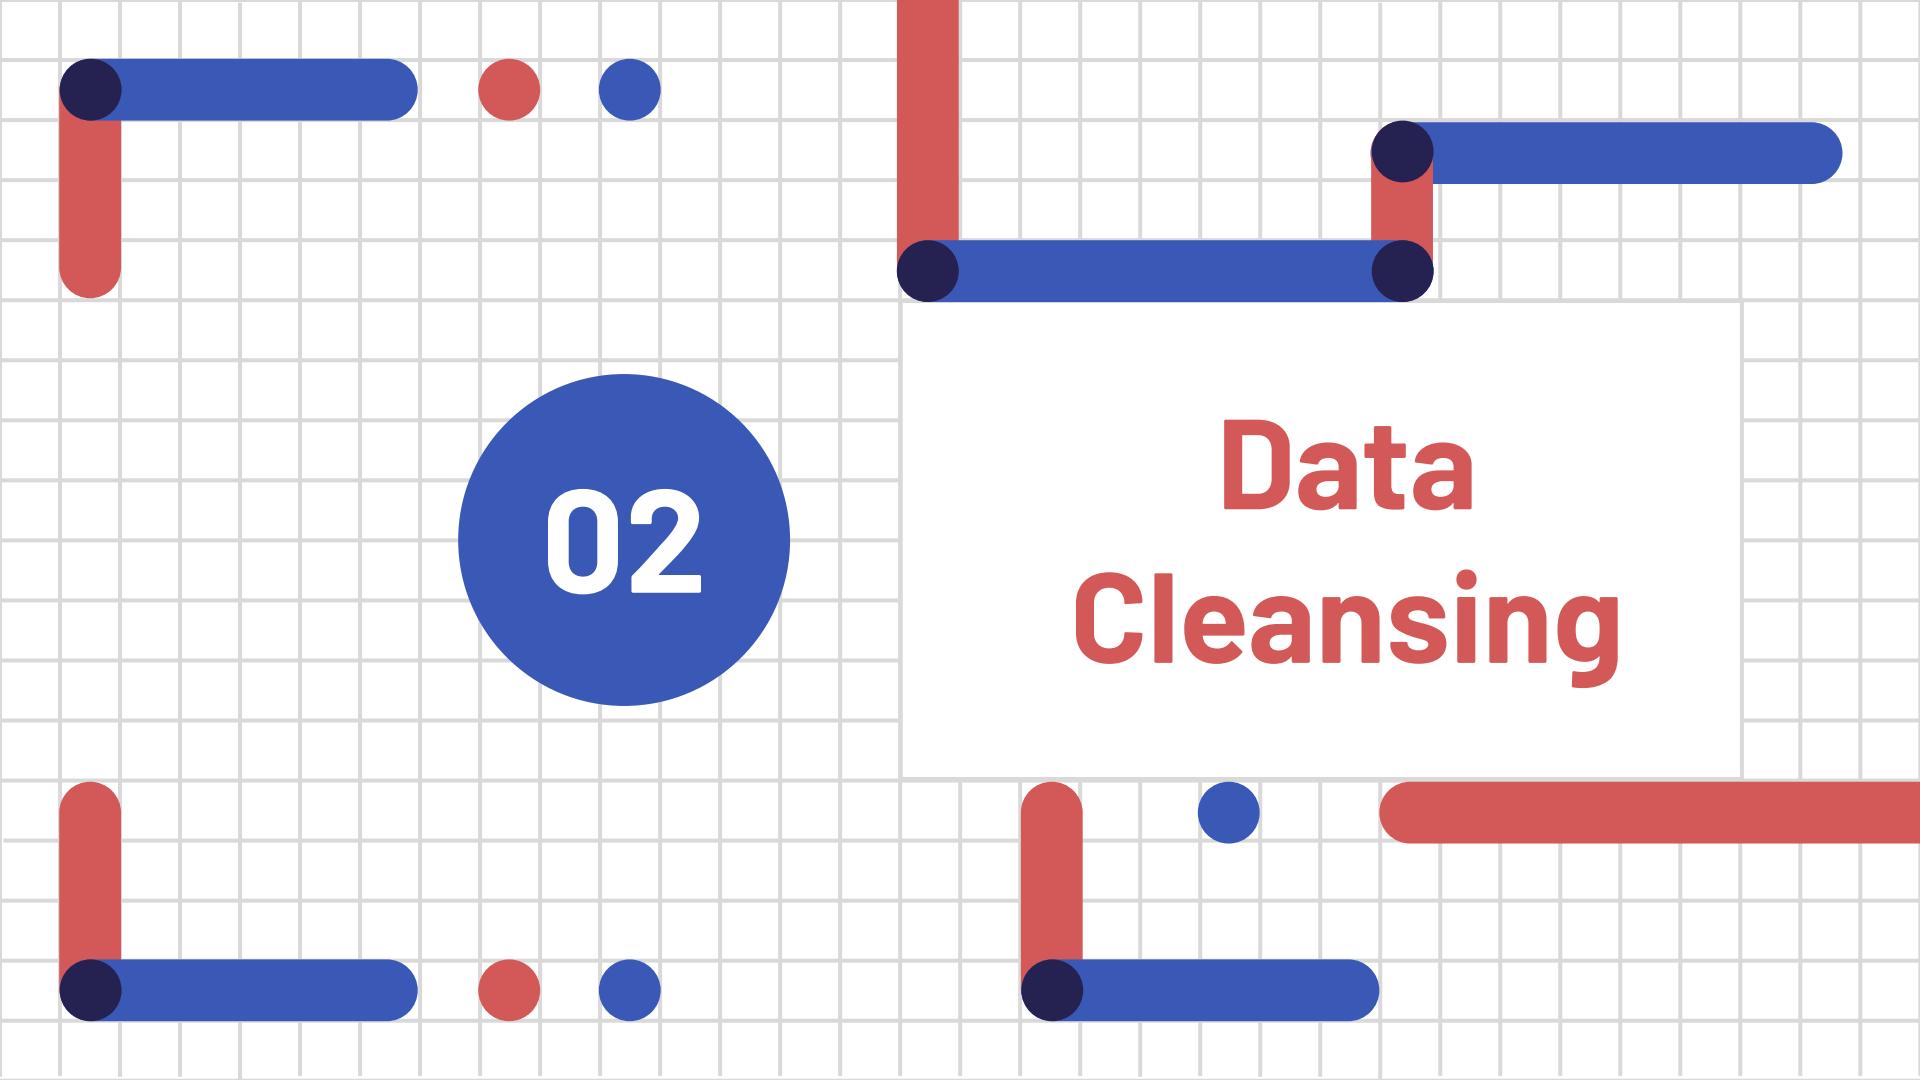
01

# Data Collected



# YouTube





02

# Data Cleansing

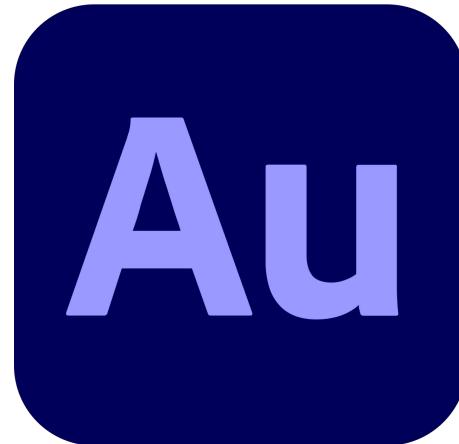


Windows 10

circleci passing | python 3.6 | 3.7 | pypi package 1.5.2 | conda-forge v1.5.2

JOSS Under Review

- Vocals (singing voice) / accompaniment separation ([2 stems](#))
- Vocals / drums / bass / other separation ([4 stems](#))
- Vocals / drums / bass / piano / other separation ([5 stems](#))





1000 files





03

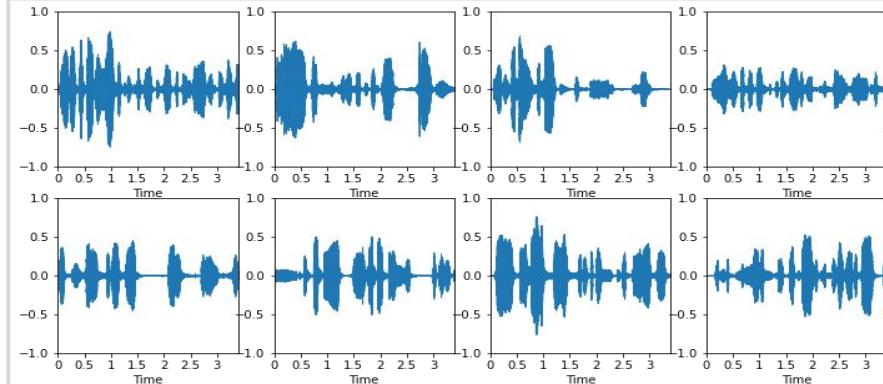
# Data Preparation

# Sound

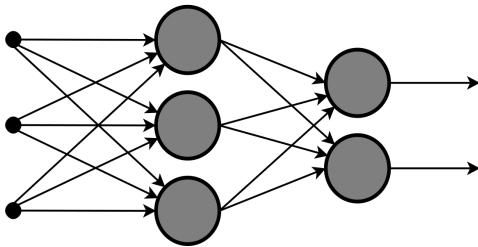
```
dataset = data_path + "sound/"  
# os.listdir(dataset)  
  
import librosa  
import librosa.display  
  
th = dataset + "TH_* (*.wav"  
cn = dataset + "CN_* (*.wav"  
en = dataset + "EN_* (*.wav"  
jp = dataset + "JP_* (*.wav"  
kr = dataset + "KR_* (*.wav"
```



```
th_signals = [librosa.load(p)[0] for p in Path().glob(th)]  
cn_signals = [librosa.load(p)[0] for p in Path().glob(cn)]  
en_signals = [librosa.load(p)[0] for p in Path().glob(en)]  
jp_signals = [librosa.load(p)[0] for p in Path().glob(jp)]  
kr_signals = [librosa.load(p)[0] for p in Path().glob(kr)]
```

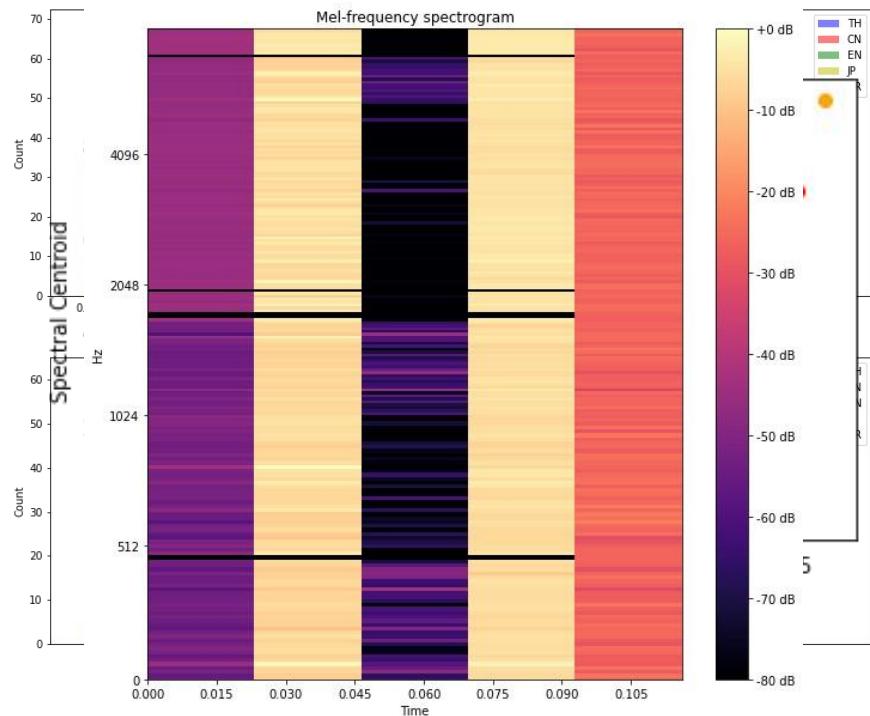


# Feature Extraction



```
def extract_features(signal):
    return [
        librosa.feature.zero_crossing_rate(signal)[0, 0],
        librosa.feature.spectral_centroid(signal)[0, 0],
        librosa.feature.mfcc(signal)[0, 0],
        librosa.feature.rms(signal)[0, 0],
        librosa.feature.spectral_contrast(signal)[0, 0],
    ]

th_features = np.array([extract_features(x) for x in th_signals])
cn_features = np.array([extract_features(x) for x in cn_signals])
en_features = np.array([extract_features(x) for x in en_signals])
jp_features = np.array([extract_features(x) for x in jp_signals])
kr_features = np.array([extract_features(x) for x in kr_signals])
```



# Split Train Test data

```
1 train_th, test_th, yth_train, yth_test = train_test_split(th_features, yth, test_size=0.2, shuffle=True,random_state=123)
2 train_cn, test_cn, ycn_train, ycn_test = train_test_split(cn_features, ycn, test_size=0.2, shuffle=True,random_state=123)
3 train_en, test_en, yen_train, yen_test = train_test_split(en_features, yen, test_size=0.2, shuffle=True,random_state=123)
4 train_jp, test_jp, yjp_train, yjp_test = train_test_split(jp_features, yjp, test_size=0.2, shuffle=True,random_state=123)
5 train_kr, test_kr, ykr_train, ykr_test = train_test_split(kr_features, ykr, test_size=0.2, shuffle=True,random_state=123)
```

```
1 train_data = np.vstack((train_th,train_cn,train_en,train_jp,train_kr))
2 y_train = yth_train+ycn_train+yen_train+yjp_train+ykr_train
3 y_train = to_categorical(y_train)
4 print("Train dataset :",train_data.shape)
5 print("Y-Train dataset :",y_train.shape)
```

Train dataset : (800, 5)  
Y-Train dataset : (800, 5)

```
1 test_data = np.vstack((test_th,test_cn,test_en,test_jp,test_kr))
2 y_test = yth_test+ycn_test+yen_test+yjp_test+ykr_test
3 y_test = to_categorical(y_test)
4 print("Test dataset :",test_data.shape)
5 print("Y-Test dataset :",y_test.shape)
```

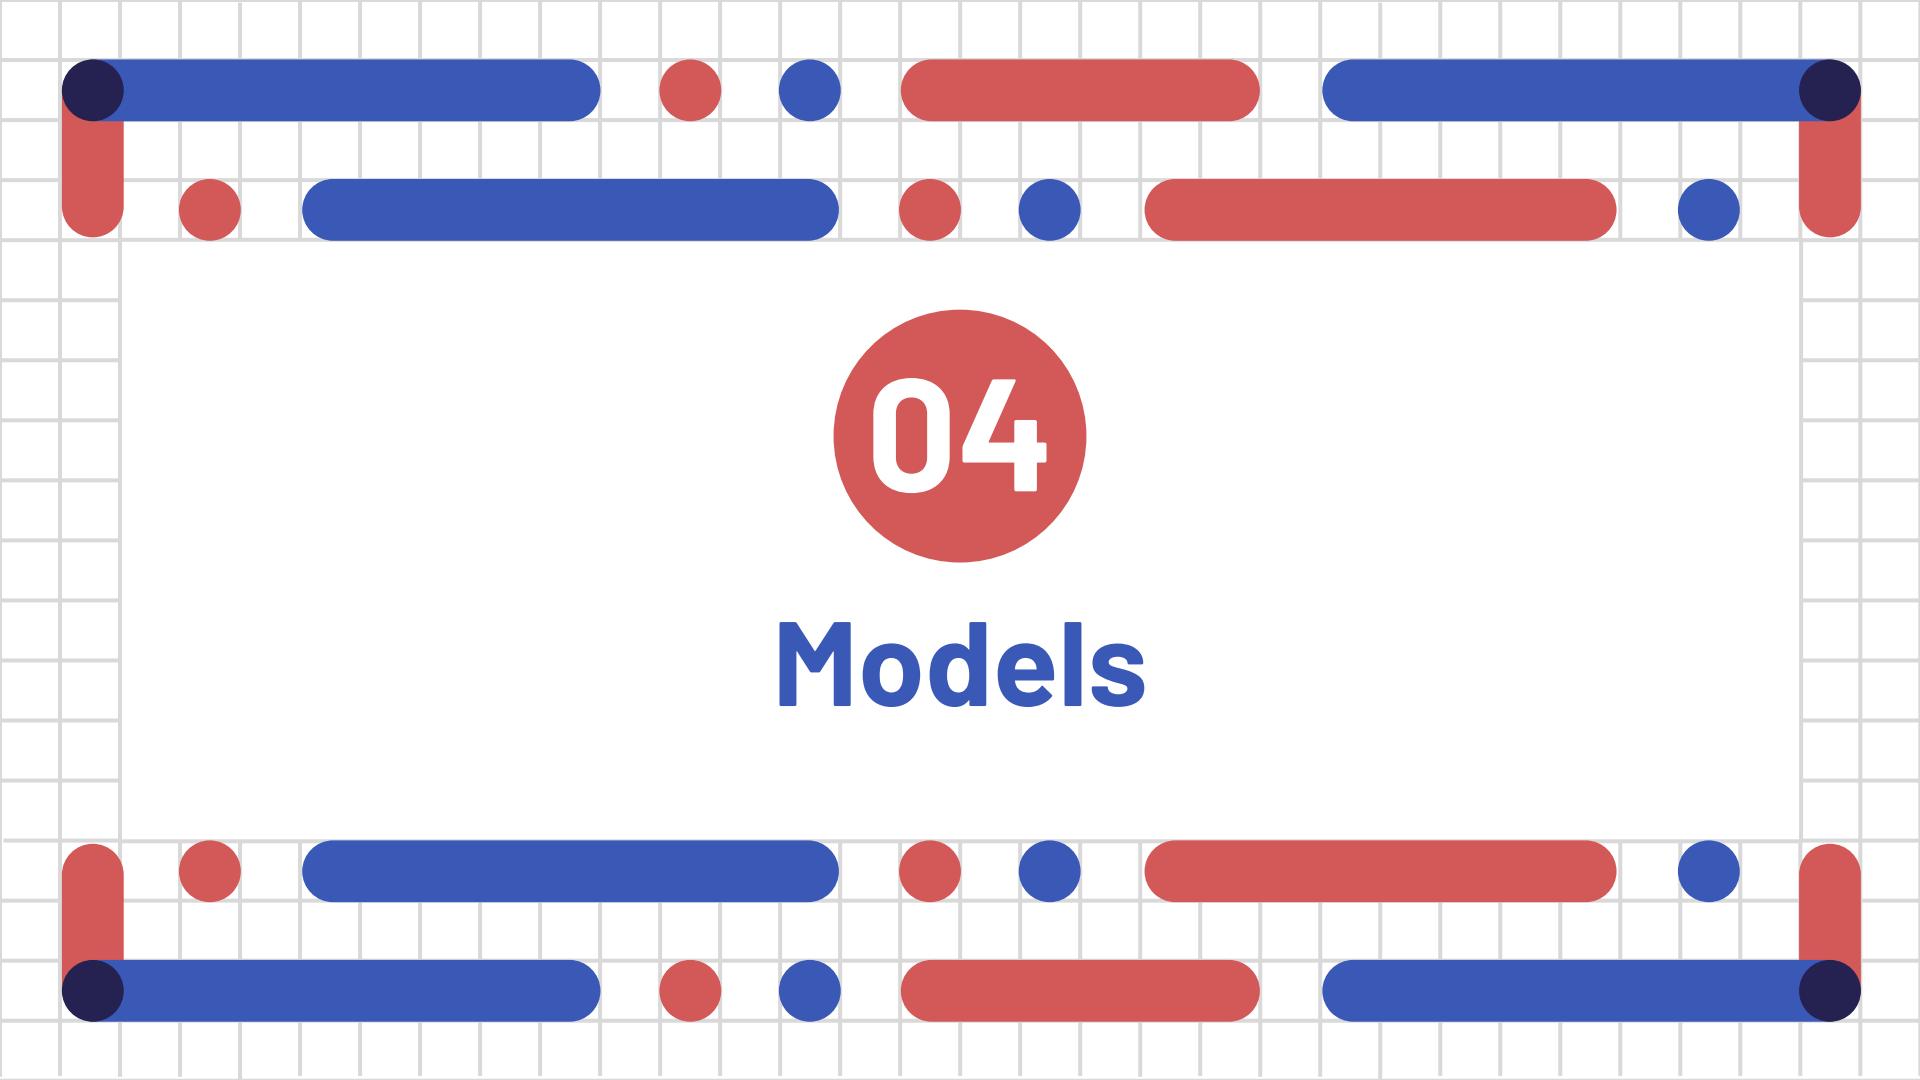
Test dataset : (200, 5)  
Y-Test dataset : (200, 5)

# Split Train Valid data

```
1 x_train, x_val, y_train, y_val = train_test_split(train_data, y_train, test_size=0.2, shuffle=True,random_state=123)

1 x_train.shape, x_val.shape, y_train.shape, y_val.shape
((640, 5), (160, 5), (640, 5), (160, 5))
```





04

# Models

# Baseline Models

```
from tensorflow.keras.models import Model
from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Activation
from tensorflow.keras.layers import LeakyReLU
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam

def create_model():
    model = Sequential()
    model.add(Dense(8, input_shape=(5,)))
    model.add(LeakyReLU(alpha=0.1))
    model.add(Dropout(0.1))
    model.add(Dense(5, activation='softmax'))
    return model

model = create_model()
model.summary()
adam_optim = Adam(learning_rate = 0.0001)
model.compile(optimizer=adam_optim, loss='categorical_crossentropy', metrics=['accuracy'])
```

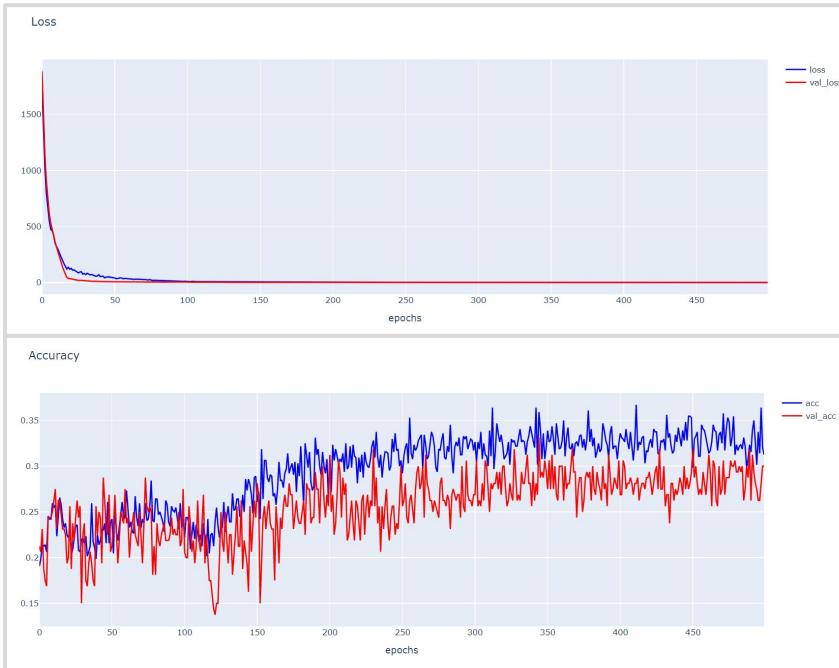
Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
dense (Dense)	(None, 8)	48
leaky_re_lu (LeakyReLU)	(None, 8)	0
dropout (Dropout)	(None, 8)	0
dense_1 (Dense)	(None, 5)	45
<hr/>		
Total params:	93	
Trainable params:	93	
Non-trainable params:	0	

# Baseline Models

```
1 his = model.fit(  
2     x_train, y_train,  
3     batch_size=1,  
4     epochs=500,  
5     verbose=1,  
6     validation_data = (x_val, y_val)  
7 )  
  
Epoch 1/500  
640/640 [=====] - 5s 4ms/step - loss: 1733.1400 - accuracy: 0.1906 - val_loss: 1885.6619 - val_accuracy: 0.2125  
Epoch 2/500  
640/640 [=====] - 2s 3ms/step - loss: 1363.1937 - accuracy: 0.1984 - val_loss: 1467.6289 - val_accuracy: 0.2062  
Epoch 3/500  
640/640 [=====] - 2s 3ms/step - loss: 1020.6646 - accuracy: 0.2141 - val_loss: 1134.4502 - val_accuracy: 0.2313  
Epoch 4/500  
640/640 [=====] - 2s 3ms/step - loss: 792.6118 - accuracy: 0.2125 - val_loss: 900.8021 - val_accuracy: 0.1875  
Epoch 5/500  
640/640 [=====] - 2s 3ms/step - loss: 674.7211 - accuracy: 0.2141 - val_loss: 738.8348 - val_accuracy: 0.1750  
Epoch 6/500  
640/640 [=====] - 2s 3ms/step - loss: 549.5569 - accuracy: 0.2062 - val_loss: 611.3173 - val_accuracy: 0.1688  
Epoch 7/500  
640/640 [=====] - 2s 3ms/step - loss: 475.3826 - accuracy: 0.2453 - val_loss: 530.2086 - val_accuracy: 0.2438  
Epoch 8/500  
640/640 [=====] - 3s 4ms/step - loss: 460.6410 - accuracy: 0.2422 - val_loss: 470.4824 - val_accuracy: 0.2438  
Epoch 9/500  
640/640 [=====] - 2s 3ms/step - loss: 420.1028 - accuracy: 0.2422 - val_loss: 417.8871 - val_accuracy: 0.2438  
Epoch 10/500  
640/640 [=====] - 2s 3ms/step - loss: 349.1039 - accuracy: 0.2594 - val_loss: 363.0685 - val_accuracy: 0.2500  
Epoch 11/500  
640/640 [=====] - 2s 4ms/step - loss: 321.4443 - accuracy: 0.2547 - val_loss: 315.3325 - val_accuracy: 0.2625
```

# Baseline Models



# Baseline Models

```
1 report = classification_report(y_true, predicted_classes, target_names=CATEGORIES, digits=4)
2
3 print(report)
```

	precision	recall	f1-score	support
TH	0.3012	0.6250	0.4065	40
CN	0.6000	0.1500	0.2400	40
EN	0.4286	0.3000	0.3529	40
JP	0.3562	0.6500	0.4602	40
KR	0.6667	0.1000	0.1739	40
accuracy			0.3650	200
macro avg	0.4705	0.3650	0.3267	200
weighted avg	0.4705	0.3650	0.3267	200

# CNN Models

```
from tensorflow.keras.models import Model
from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Activation
from tensorflow.keras.layers import LeakyReLU
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Conv2D

from tensorflow.keras.layers import GlobalAveragePooling2D
from tensorflow.keras.layers import MaxPooling2D
from tensorflow.keras.optimizers import Adam

file_name = 'drive/My Drive/AUCC/DP/dataset/sound/TH_M (1).wav'
audio, sample_rate = librosa.load(file_name, res_type='kaiser_fast')

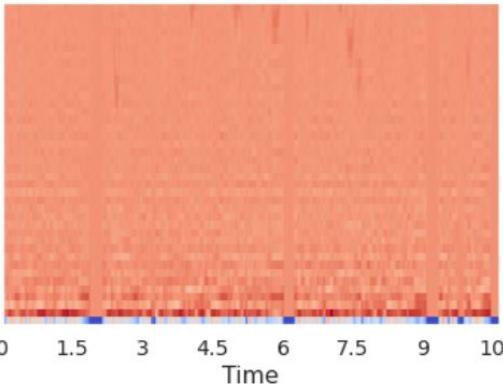
print(audio.shape, sample_rate)

mfccs = librosa.feature.mfcc(y=audio, sr=sample_rate, n_mfcc=40)
print(mfccs.shape)

(233856,) 22050
(40, 457)
```

```
1   librosa.display.specshow(mfccs, sr=sample_rate, x_axis='time')
```

```
<matplotlib.collections.QuadMesh at 0x7fc64d3ab9d0>
```



# CNN Models

```
1 max_pad_len = 1000
2
3 def extract_features(file_name):
4     audio, sample_rate = librosa.load(file_name, res_type='kaiser_fast')
5     mfccs = librosa.feature.mfcc(y=audio, sr=sample_rate, n_mfcc=40)
6     pad_width = max_pad_len - mfccs.shape[1]
7     mfccs = np.pad(mfccs, pad_width=((0, 0), (0, pad_width)), mode='constant')
8     return mfccs
```

```
1 TH_features = []
2 for file_name in Path().glob(th):
3     data = extract_features(file_name)
4     TH_features.append(data)
5 print(len(TH_features))
6 TH_features[0].shape
```

```
200
(40, 1000)
```

```
1 CN_features = []
2 for file_name in Path().glob(cn):
3     data = extract_features(file_name)
4     CN_features.append(data)
5 print(len(CN_features))
6 CN_features[0].shape
```

```
200
(40, 1000)
```

```
1 EN_features = []
2 for file_name in Path().glob(en):
3     data = extract_features(file_name)
4     EN_features.append(data)
5 print(len(EN_features))
6 EN_features[0].shape
```

```
200
(40, 1000)
```

```
1 JP_features = []
2 for file_name in Path().glob(jp):
3     data = extract_features(file_name)
4     JP_features.append(data)
5 print(len(JP_features))
6 JP_features[0].shape
```

```
200
(40, 1000)
```

```
1 KR_features = []
2 for file_name in Path().glob(kr):
3     data = extract_features(file_name)
4     KR_features.append(data)
5 print(len(KR_features))
6 KR_features[0].shape
```

```
200
(40, 1000)
```

# CNN Models

```
num_rows = 40
num_columns = 1000
num_channels = 1
num_batch_size = 64
epochs = 500

model = Sequential()
model.add(Conv2D(filters=16, kernel_size=2, input_shape=(num_rows, num_columns, num_channels), activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.2))

model.add(Conv2D(filters=32, kernel_size=2, activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.2))

model.add(Conv2D(filters=64, kernel_size=2, activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.2))

model.add(Conv2D(filters=128, kernel_size=2, activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Dropout(0.2))
model.add(GlobalAveragePooling2D())

model.add(Dense(5, activation='softmax'))

model.summary()
model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 39, 999, 16)	80
max_pooling2d (MaxPooling2D)	(None, 19, 499, 16)	0
dropout (Dropout)	(None, 19, 499, 16)	0
conv2d_1 (Conv2D)	(None, 18, 498, 32)	2080
max_pooling2d_1 (MaxPooling2D)	(None, 9, 249, 32)	0
dropout_1 (Dropout)	(None, 9, 249, 32)	0
conv2d_2 (Conv2D)	(None, 8, 248, 64)	8256
max_pooling2d_2 (MaxPooling2D)	(None, 4, 124, 64)	0
dropout_2 (Dropout)	(None, 4, 124, 64)	0
conv2d_3 (Conv2D)	(None, 3, 123, 128)	32896
max_pooling2d_3 (MaxPooling2D)	(None, 1, 61, 128)	0
dropout_3 (Dropout)	(None, 1, 61, 128)	0
global_average_pooling2d (GlobalAveragePooling2D)	(None, 128)	0
dense (Dense)	(None, 5)	645

Total params: 43,957

Trainable params: 43,957

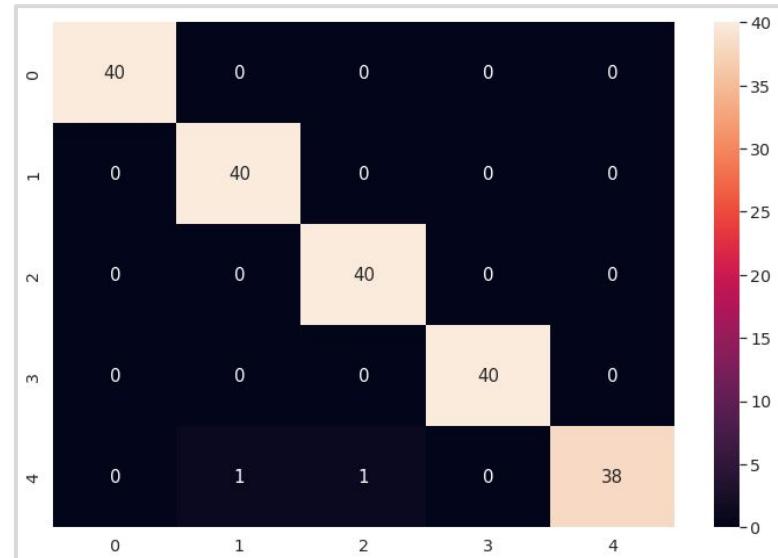
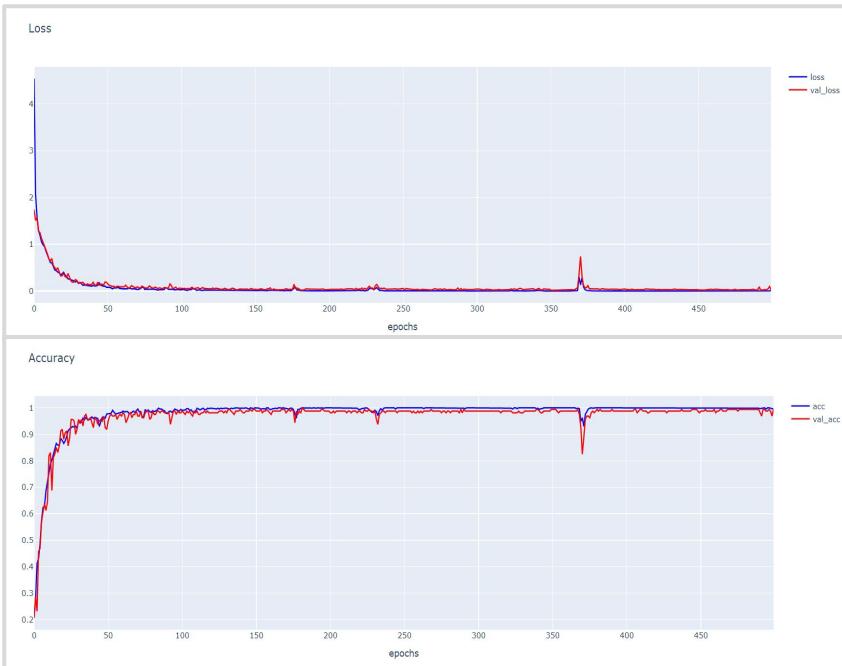
Non-trainable params: 0

# CNN Models

```
1 his = model.fit(x_train, y_train,
2                   batch_size=num_batch_size,
3                   epochs=epochs,
4                   validation_data=(x_val, y_val),
5                   verbose=1)

Epoch 1/500
10/10 [=====] - 1s 133ms/step - loss: 4.5420 - accuracy: 0.2078 - val_loss: 1.7414 - val_accuracy: 0.2062
Epoch 2/500
10/10 [=====] - 1s 89ms/step - loss: 2.0627 - accuracy: 0.2766 - val_loss: 1.5087 - val_accuracy: 0.2875
Epoch 3/500
10/10 [=====] - 1s 90ms/step - loss: 1.5829 - accuracy: 0.4125 - val_loss: 1.5723 - val_accuracy: 0.2313
Epoch 4/500
10/10 [=====] - 1s 79ms/step - loss: 1.2940 - accuracy: 0.4297 - val_loss: 1.2806 - val_accuracy: 0.4563
Epoch 5/500
10/10 [=====] - 1s 80ms/step - loss: 1.1900 - accuracy: 0.4859 - val_loss: 1.2474 - val_accuracy: 0.4688
Epoch 6/500
10/10 [=====] - 1s 79ms/step - loss: 1.0540 - accuracy: 0.5734 - val_loss: 1.1262 - val_accuracy: 0.5688
Epoch 7/500
10/10 [=====] - 1s 80ms/step - loss: 0.9888 - accuracy: 0.6234 - val_loss: 1.0489 - val_accuracy: 0.6062
Epoch 8/500
10/10 [=====] - 1s 90ms/step - loss: 0.9526 - accuracy: 0.6281 - val_loss: 0.9537 - val_accuracy: 0.6375
Epoch 9/500
10/10 [=====] - 1s 80ms/step - loss: 0.8550 - accuracy: 0.6859 - val_loss: 0.8805 - val_accuracy: 0.6125
Epoch 10/500
10/10 [=====] - 1s 81ms/step - loss: 0.7669 - accuracy: 0.7188 - val_loss: 0.7960 - val_accuracy: 0.6438
Epoch 11/500
10/10 [=====] - 1s 82ms/step - loss: 0.7065 - accuracy: 0.7453 - val_loss: 0.6872 - val_accuracy: 0.8188
Epoch 12/500
10/10 [=====] - 1s 80ms/step - loss: 0.6102 - accuracy: 0.7875 - val_loss: 0.6426 - val_accuracy: 0.8313
```

# CNN Models



# CNN Models

```
1 report = classification_report(y_true, predicted_classes, target_names=CATEGORIES, digits=4)
2
3 print(report)
```

	precision	recall	f1-score	support
TH	1.0000	1.0000	1.0000	40
CN	0.9756	1.0000	0.9877	40
EN	0.9756	1.0000	0.9877	40
JP	1.0000	1.0000	1.0000	40
KR	1.0000	0.9500	0.9744	40
accuracy			0.9900	200
macro avg	0.9902	0.9900	0.9899	200
weighted avg	0.9902	0.9900	0.9899	200

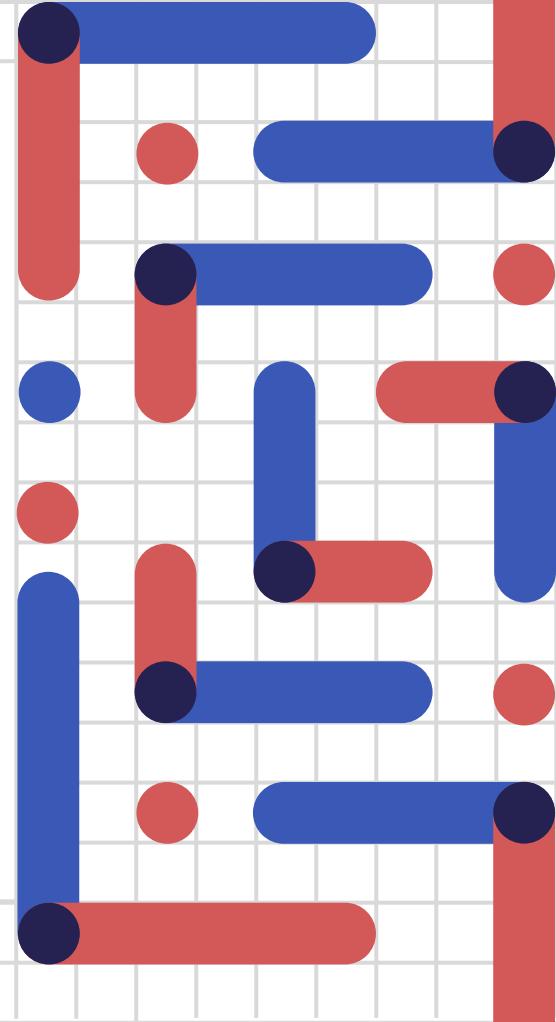
# Thanks

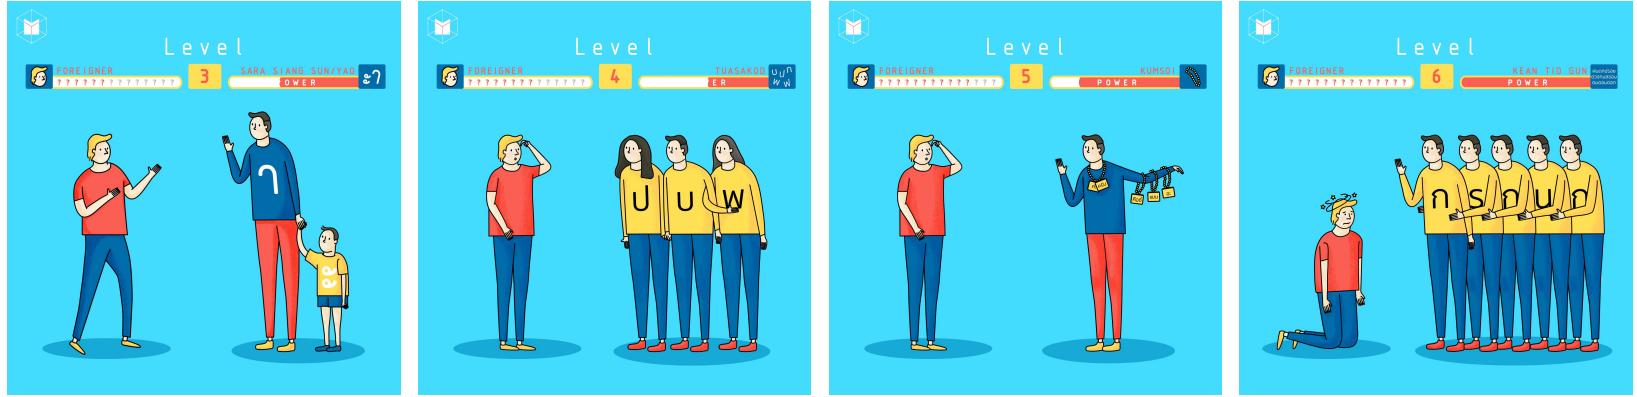
620710405 นางสาวณัฐริดา ลากานชัย

620710745 นางสาวอาทิตย์รา ชุมทอง



CREDITS: This presentation template was created by [Slidesgo](#), including icons by [Flaticon](#), infographics & images by [Freepik](#)





<https://thematter.co/social/thai-language-is-not-that-easy/29010>